

NIRVANA – Master Specification

Volume 8: AI Advisory Engine

Purpose

The AI layer explains, personalizes and communicates recommendations. It never replaces deterministic financial calculations.

Architecture

User Query → Context Builder → Recommendation Engine Output → AI Prompt Assembly → LLM → Validation → Response.

System Rules

The LLM must only use normalized financial data supplied by Nirvana. It should never invent returns, prices or tax rules.

Capabilities

Explain recommendations, answer financial questions, compare scenarios, summarize portfolios, coach behavioural biases, generate client-ready reports and answer follow-up questions.

Prompt Strategy

Provide the model with user profile, goals, portfolio snapshot, market context and recommendation outputs before generating explanations.

Safety

Flag uncertainty, distinguish assumptions from facts, recommend professional advice where appropriate and maintain audit logs of prompts and outputs.

Report Generation

Generate Wealth Health Reports, Retirement Reviews, Tax Summaries, Goal Progress Reports, Withdrawal Plans and Human vs AI comparison reports in PDF-ready format.